

# MPP-E1180 Lecture 3: Introduction to the R Programming Language

Christopher Gandrud

25 September 2014

# Objectives for the week

- ▶ Review
- ▶ Sync a fork with the original
- ▶ Reminder: **Pair Assignment 1**
- ▶ Basics of object-oriented programming in R
- ▶ Simple R data structures
- ▶ Simple descriptive statistics and plotting with Base R

# Class make up

We need to reschedule the classes originally on **30 October** and **13 November**.

Possibilities:

- ▶ Sometime in Midterm Week
- ▶ Mondays, 14-16h & 18-20h
- ▶ Tuesdays, 8-10h & 18-20h
- ▶ Wednesdays, 8-10h & 18-20h
- ▶ Thursdays, 14-16h (double session)

# Review

With a partner:

- ▶ What is the difference between **relative** and **absolute** file paths?
- ▶ What is a **commit**?
- ▶ What does it mean to **pull** and **push** a repo?

## Sync a fork with the original repo

See: <https://help.github.com/articles/syncing-a-fork>.

Change current working directory to the repo:

```
cd /FILE/PATH
```

Add the original repo as an **upstream repo**:

# Note: should be on one line & only need to do once

```
git remote add upstream
```

```
https://github.com/HertieDataScience2014/SyllabusAndLecture
```

Then **fetch** and **merge**:

```
git fetch upstream
```

```
git checkout master
```

```
git merge upstream/master
```

# Pair Assignment 1

- ▶ **Due:** Midnight 26 September.
- ▶ Learning objectives: develop your understanding of
  - ▶ **file structures,**
  - ▶ **version control,**
  - ▶ **basic R data structures** and **descriptive statistics.**

# Pair Assignment 1

Each pair will create a **new public GitHub repository**

- ▶ Must be **fully documented**, including with a **descriptive README.md** file. Your code must be **human readable** and **clearly commented**.
- ▶ Include **R source** code files that:
  - ▶ Access at least **two** core R data sets
  - ▶ Illustrate the datas' distributions using a variety of **relevant descriptive statistics**
  - ▶ Two files must be **dynamically linked**
- ▶ **Another pair** makes a **pull request**. And this is discussed/merged.

# What is R?

**Open source programming language**, with a particular focus on statistical programming.

**History:** Originally (in 1993) an implementation of the S programming language (Bell Labs), by **R**oss Ihaka and **R**obert Gentleman (hence **R**) at University of Auckland.

Currently the R Foundation for Statistical Computing is based in Vienna.

**RStudio:**

R is to RStudio as Git is to GitHub.



## Growing popularity

R can be easily expanded by **user created packages** hosted on GitHub and/or CRAN.

# How to Cite R

```
citation()
```

```
##
```

```
## To cite R in publications use:
```

```
##
```

```
## R Core Team (2014). R: A language and environment for  
## statistical computing. R Foundation for Statistical Computing  
## Vienna, Austria. URL http://www.R-project.org/.
```

```
##
```

```
## A BibTeX entry for LaTeX users is
```

```
##
```

```
## @Manual{,
```

```
## title = {R: A Language and Environment for Statistical Computing},
```

```
## author = {{R Core Team}},
```

```
## organization = {R Foundation for Statistical Computing},
```

```
## address = {Vienna, Austria},
```

```
## year = {2014},
```

```
## url = {http://www.R-project.org/},
```

# Fundamentals of the R language

R is **object-oriented**.

**Objects are R's nouns.** They include (not exhaustive):

- ▶ character string (e.g. words)
- ▶ number
- ▶ vector of numbers or character strings
- ▶ matrix
- ▶ data frame
- ▶ list

# Assignment

You use the **assignment operator** (<-) to assign character strings, numbers, vectors, etc. to object names

```
## Assign the number 10 to an object called number  
number <- 10
```

```
number
```

```
## [1] 10
```

```
# Assign Hello world to an object called words  
words <- "Hello World"
```

```
words
```

```
## [1] "Hello World"
```

# Assignment

You can also use =:

```
number = 10
```

```
number
```

```
## [1] 10
```

Note: it has a slightly different meaning.

See StackOverflow discussion.

# Special values in R

- ▶ NA: not available, missing
- ▶ NULL: does not exist, is undefined
- ▶ TRUE, T: logical true. **Logical** is also an object class.
- ▶ FALSE, F: logical false

# Finding special values

Function	Meaning
<code>is.na</code>	Is the value NA
<code>is.null</code>	Is the value NULL
<code>isTRUE</code>	Is the value TRUE
<code>!isTRUE</code>	Is the value FALSE

```
absent <- NA  
is.na(absent)
```

```
## [1] TRUE
```

---

Operator	Meaning
<	less than
>	greater than
==	equal to
<=	less than or equal to
>=	greater than or equal to
!=	not equal to
a   b	a or b
a & b	a and b

---



# Classes

Objects have distinct classes.

```
# Find the class of number  
class(number)
```

```
## [1] "numeric"
```

```
# Find the class of absent  
class(absent)
```

```
## [1] "logical"
```

# Naming objects

- ▶ Object names **cannot have spaces**
  - ▶ Use CamelCase, name\_underscore, or name.period
- ▶ Avoid creating an object with the same name as a function (e.g. `c` and `t`) or special value (`NA`, `NULL`, `TRUE`, `FALSE`).
- ▶ Use **descriptive object names!**
  - ▶ Not: `obj1`, `obj2`
- ▶ Each object name must be **unique** in an environment.
  - ▶ Assigning something to an object name that is already in use will **overwrite the object's previous contents**.

# Finding objects

```
# Find objects in your workspace  
ls()
```

```
## [1] "absent" "number" "words"
```

# Style Guides

As with natural language writing, it is a good idea to stick to one style guide with your R code:

- ▶ Google's R Style Guide
- ▶ Hadely Wickham's R Style Guide

# Vectors

A vector is an **ordered collection** of numbers, characters, etc. of the **same type**.

Vectors can be created with the `c` (**combine**) function.

```
# Create numeric vector
```

```
numeric_vector <- c(1, 2, 3)
```

```
# Create character vector
```

```
character_vector <- c('Albania', 'Botswana', 'Cambodia')
```

## Factor class vector

Categorical variables are called **factors** in R.

```
# Create numeric vector
fruits <- c(1, 1, 2)

# Create character vector for factor labels
fruit_names <- c('apples', 'mangos')

# Convert to labelled factor
fruits_factor <- factor(fruits, labels = fruit_names)

summary(fruits_factor)
```

```
## apples mangos
##      2      1
```

# Matrices

Matrices are collections of vectors **with the same length**

```
# Combine numeric_vector and character_vector into a matrix  
combined <- cbind(numeric_vector, character_vector)
```

```
combined
```

```
##      numeric_vector character_vector  
## [1,] "1"           "Albania"  
## [2,] "2"           "Botswana"  
## [3,] "3"           "Cambodia"
```

Note: In addition to `cbind` you can `rbind` new rows onto a matrix.

# Data frames

Data frames are collections of vectors with the same length.

Each column (vector) can be of a **different class**.

```
# Combine numeric_vector and character_vector into a data frame  
combined_df <- data.frame(numeric_vector, character_vector,  
                           stringsAsFactors = FALSE)
```

```
combined_df
```

```
##   numeric_vector character_vector  
## 1              1           Albania  
## 2              2           Botswana  
## 3              3           Cambodia
```



# Lists

A list is a vector containing other objects.

The objects can have different lengths and classes.

```
# Create a list with three objects of different lengths
test_list <- list(countries = character_vector, not_there =
                  more_numbers = 1:100)
```

```
test_list
```

```
## $countries
```

```
## [1] "Albania" "Botswana" "Cambodia"
```

```
##
```

```
## $not_there
```

```
## [1] NA NA
```

```
##
```

```
## $more_numbers
```

```
## [1] 1 2 3 4 5 6 7 8 9 10 11 12 13
```

```
## [18] 18 19 20 21 22 23 24 25 26 27 28 29 30
```

# Functions

Functions do things to/with objects. Functions are like **R's verbs**.

When using functions to do things to objects, they are always followed by parentheses (). The parentheses contain the **arguments**. Arguments are separated by commas.

```
# Summarise combined_df  
summary(combined_df, digits = 2)
```

```
##  numeric_vector character_vector  
##  Min.      :1.0      Length:3  
##  1st Qu.:1.5      Class :character  
##  Median :2.0      Mode  :character  
##  Mean     :2.0  
##  3rd Qu.:2.5  
##  Max.     :3.0
```

# Functions help

Use ? to find out what arguments a function can take.

```
?summary
```

The help page will also show the function's **default argument values**.

## Component selection (\$)

The \$ is known as the component selector. It selects a component of an object.

```
combined_df$character_vector
```

```
## [1] "Albania" "Botswana" "Cambodia"
```

## Subscripts []

You can use subscripts [] to also select components.

For data frames they have a [row, column] pattern.

```
# Select the second row and first column of combined_df  
combined_df[2, 1]
```

```
## [1] 2
```

```
# Select the first two rows  
combined_df[c(1, 2), ]
```

```
##      numeric_vector character_vector  
## 1             1           Albania  
## 2             2           Botswana
```

## Subscripts []

```
# Select the character_vector column  
combined_df[, 'character_vector']
```

```
## [1] "Albania" "Botswana" "Cambodia"
```

## Assignment with elements of objects

You can use assignment with parts of objects. For example:

```
combined_df$character_vector[3] <- 'China'  
combined_df$character_vector
```

```
## [1] "Albania" "Botswana" "China"
```

You can even add new variables:

```
combined_df$new_var <- 1:3  
combined_df
```

```
##   numeric_vector character_vector new_var  
## 1             1           Albania      1  
## 2             2           Botswana      2  
## 3             3             China      3
```

# Packages

You can greatly expand the number of functions available to you by installing and loading user created packages.

```
# Install dplyr package  
install.packages('dplyr')
```

```
# Load dplyr package  
library(dplyr)
```

You can also call a function directly from a specific package with the double colon operator (`::`).

```
Grouped <- dplyr::group_by(combined_df, character_vector)
```



# R's build-in data sets

List internal data sets:

```
data()
```

Load **swiss** data set:

```
data(swiss)
```

Find data description:

```
?swiss
```

## R's build-in data sets

Find variable names:

```
names(swiss)
```

```
## [1] "Fertility"      "Agriculture"    "Examination"  
## [4] "Education"     "Catholic"       "Infant.Mortal"
```

See the first three rows and four columns

```
head(swiss[1:3, 1:4])
```

##	Fertility	Agriculture	Examination	Education
## Courtelary	80.2	17.0	15	12
## Delemont	83.1	45.1	6	9
## Franches-Mnt	92.5	39.7	5	5

# What all the cool kids are doing: piping

**Pipe:** pass a value forward to a function call.

Why?

- ▶ Faster compilation.
- ▶ Enhanced code readability.

In R use `%>%` from the `magrittr` package.

`%>%` passes a value to the **first argument** of the next function call.

## Simple piping example

Not piped:

```
values <- rnorm(1000, mean = 10)
value_mean <- mean(values)
round(value_mean, digits = 2)
```

```
## [1] 9.98
```

Piped:

```
library(magrittr)

rnorm(1000, mean = 10) %>% mean() %>% round(digits = 2)
```

```
## [1] 9.99
```

# Descriptive statistics: review


**Descriptive Statistics:** describe samples

Stats 101: describe samples **distributions** with appropriate measure of

- ▶ **central tendency**
- ▶ **variability**


# Histograms

```
hist(swiss$Examination)
```



# Histograms: styling

```
hist(swiss$Examination,  
     main = 'Swiss Canton Draftee Examination Scores (1888)',  
     xlab = '% receiving highest mark on army exam')
```



## Digression: Creating functions

You can create a function to find the sample mean ( $\bar{x} = \frac{\sum x}{n}$ ) of a vector.

```
fun_mean <- function(x){  
  sum(x) / length(x)  
}  
  
## Find the mean  
fun_mean(x = swiss$Examination)  
  
## [1] 16.49
```



# Finding means

(or use the mean function in base R)

```
mean(swiss$Examination)
```

```
## [1] 16.49
```

If you have missing values (NA):

```
mean(swiss$Examination, na.rm = TRUE)
```

## Digression: Loops

You can 'loop' through the data set to find the mean for each column

```
for (i in 1:length(names(swiss))) {  
  swiss[, i] %>%  
    mean() %>%  
    round(digits = 1) %>%  
    paste(names(swiss)[i], ., '\n') %>% # the . directs the  
    cat()  
}
```

```
## Fertility 70.1  
## Agriculture 50.7  
## Examination 16.5  
## Education 11  
## Catholic 41.1  
## Infant.Mortality 19.9
```

# Other functions for central tendency

## Median

```
median(swiss$Examination)
```

```
## [1] 16
```


## Mode

mode is not an R function to find the statistical mode.

Instead use `summary` for factor nominal variables or make a bar chart.

## Simple bar chart for nominal

```
devtools::source_url('http://bit.ly/OTWEGS')  
plot(MortalityGDP$region, xlab = 'Region')
```



# Variation

## Range:

```
range(swiss$Examination)
```

```
## [1] 3 37
```

## Quartiles:


```
summary(swiss$Examination)
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
##	3.0	12.0	16.0	16.5	22.0	37.0

# Variation

## Boxplots:

```
boxplot(swiss$Examination, main = '% of Draftees with Highest Mark')
```



# Variation

**Interquartile Range ( $IQR = Q_3 - Q_1$ ):**

```
IQR(swiss$Examination)
```

```
## [1] 10
```

## Variation: standard deviation

**Sum of squared deviations:**

$$\text{Sum of Squares} = \sum (x - \bar{x})^2$$

**Degrees of freedom** (number of values that are free to vary):

$$\text{df} = n - 1$$

**Variance ( $s^2$ ):**

$$s^2 = \frac{\text{Sum of Squares}}{\text{Degrees of Freedom}} = \frac{\sum (x - \bar{x})^2}{n - 1}$$

**Standard deviation ( $s$ )** (in terms of the mean):

$$s = \sqrt{s^2}$$



## Variation: Standard Error

The **standard error** of the mean:

If we think of the variation as around a central tendency as a measure of **unreliability** then we want the measure to **decrease as the sample size goes up**.

$$SE_{\bar{x}} = \frac{s}{\sqrt{n}}$$

# Variation: Variance and Standard Deviation

## Variance:

```
var(swiss$Examination)
```

```
## [1] 63.65
```

## Standard Deviation:

```
sd(swiss$Examination)
```

```
## [1] 7.978
```

## Variation: Standard Error

### Standard Error:

```
sd_error <- function(x) {  
  sd(x) / sqrt(length(x))  
}
```


```
sd_error(swiss$Examination)
```

```
## [1] 1.164
```

# Playing with distributions

Simulated normally distributed data with SD of 30 and mean 50

```
Normal130 <- rnorm(1e+6, mean = 50, sd = 30)
```




# Transform skewed data

Highly skewed data can be transformed to have a normal distribution.

Helps correct two violations of key assumptions: (a) non-linearity and (b) heteroskedasticity.


```
hist(swiss$Education, main = '')
```



# Natural log transformed skewed data


```
log(swiss$Education) %>% hist()
```

Histogram of log(swiss\$Education)



# Joint distributions

```
plot(log(swiss$Education), swiss$Examination)
```



## Summarise with correlation coefficients

```
cor.test(log(swiss$Education), swiss$Examination)
```

```
##
```

```
## Pearson's product-moment correlation
```

```
##
```

```
## data: log(swiss$Education) and swiss$Examination
```

```
## t = 6.431, df = 45, p-value = 7.133e-08
```

```
## alternative hypothesis: true correlation is not equal to 0
```

```
## 95 percent confidence interval:
```

```
## 0.5053 0.8169
```

```
## sample estimates:
```


```
## cor
```

```
## 0.6921
```



## Summarise with loess

```
ggplot2::ggplot(swiss, aes(log(Education), Examination)) +  
  geom_point() + geom_smooth() + theme_bw()
```



# Programming Hint (1)

***Always close!***

In R this means closing:

▶ `()`

▶ `[]`

▶ `{ }`

▶ `' '`

▶ `" "`

## Programming Hint (2)

Make your code as **simple as possible**.

- ▶ Easier to read.
- ▶ Easier to write (ultimately).
- ▶ Easier to find mistakes.
- ▶ Often computationally more efficient.

One way to do this is to **define things once**.

## Programming Hint (2)

### Bad

```
mean(rnorm(1000))
```

```
## [1] -0.02098
```

```
sd(rnorm(1000))
```

```
## [1] 0.9898
```

## Programming Hint (2)

### Good

```
rand_sample <- rnorm(1000)
```

```
mean(rand_sample)
```

```
## [1] 0.0002194
```

```
sd(rand_sample)
```

```
## [1] 1.013
```

# Seminar: Start using R!

- ▶ **Access** R data sets
- ▶ Explore the data and find ways to **numerically/graphically** describe it.
- ▶ Find and use R functions that were **not covered** in the lecture for exploring and transforming your data.
- ▶ Create **your own function** (what it does is open to you).